

Cleaning a Conversation Transcript in Praat

You have been given a folder of code and a recording. The code writes a first draft of the transcript; you then make it **line up** — the right words, on the right speaker, starting at the exact moment they were spoken — and turn it into the spreadsheet the lab analyses. This guide goes click by click.

Hayden Lab · Baylor College of Medicine · James Belanger

1 · Whisper

The code writes down what was said and guesses who said it. Good at words; its timings are late.

2 · MFA

The code then re-finds those words in the audio, which fixes the timings. Both steps are one command.

3 · You, in Praat

Fix what the machines got wrong: words, speakers, and the boundaries they could not resolve.

4 · Spreadsheet

One more command turns your finished file into the Excel table the lab analyses.

WHAT YOU ARE DOING

Open the recording next to its transcript, walk through it, and fix three things — **the words, who said them, and where each word starts**. Then save and hand it back.

Why this matters

These recordings were made while someone had electrodes in their brain, and we were watching individual neurons fire while they talked and listened. The entire experiment is the **alignment** of those two things: what the brain did, and what was being said at that exact moment. A transcript that says the right words but puts them at the wrong times is, for our purposes, no transcript at all.

The precision needed is finer than it sounds. A neuron's response to a word plays out over a few hundred milliseconds, and we measure it in bins tens of milliseconds wide. The machine's word onsets land about **65 ms late** — on the same scale as the effect we are trying to see. That error does not average away: every word is late by roughly the same amount, so the whole response shifts, a sharp burst of firing smears into a low bump, and a real effect can vanish or a false one appear. Correcting the boundaries recovers it from the very same recording. That is what this job is.

THREE THINGS YOU FIX, THREE WAYS THE ANALYSIS BREAKS WITHOUT YOU

WHAT YOU FIX	WHAT IT COSTS TO GET WRONG
the word itself	every word is looked up in models of meaning and grammar. A wrong word gets the wrong properties and is never questioned again — the error just becomes part of the data
who said it	speaking and listening drive neurons in completely different ways. A word on the wrong speaker's row does not add noise, it moves that moment into the wrong experimental condition
where the word starts	the onset is the moment everything is measured from. Off by 65 ms and you are measuring the wrong slice of the brain's response to that word

The machine is genuinely good at words and systematically wrong about time; your ear and the waveform on screen are what fix it. And nothing downstream re-checks this — from here on the analysis is automatic, so **you are the last person who looks**. Work carefully, and flag anything that seems wrong rather than guessing.

0 Which job do you have?

There are two ways this lands on your desk, and they need different amounts of this document. Check which one you are on before reading any further.

A · YOU WERE GIVEN A RECORDING

A sound file and a folder of code, and **no transcript yet**. You make the first draft yourself, then clean it.

Read the whole thing, starting at section 1.

B · YOU WERE GIVEN A TRANSCRIPT ALREADY

A sound file **and** a `.TextGrid` — someone has already run the code for you. Your job is the cleaning, which is the part that needs a human.

Skip sections 1, 2 and 3 completely and start at section 4. You will never need Python, conda or a command window.

If nobody handed you a `.TextGrid`, you are on route A. If you are not sure, ask — it takes one message and saves you an hour of setting up software you did not need.

The kinds of file you will meet

LOOKS LIKE	WHAT IT IS
<code>something.wav</code>	the audio recording itself
<code>something.TextGrid</code>	Praat's transcript file: the words, and the exact moment each one starts and ends. Opened in a text editor it looks like gibberish — always open it in Praat
<code>something.xlsx</code>	an Excel spreadsheet — what the lab analyses, and what you hand back
<code>something.py</code>	a piece of code you run by typing a command. Double-clicking it does nothing useful

1 Get the code ROUTE A ONLY

The code lives on a web page. You download it once, as a single zip file — you do not need an account, and you do not need to know what GitHub is.

1 Open the page

github.com/BCM-Neurosurgery/conversation-alignment

The screenshot shows the GitHub repository page for `BCM-Neurosurgery/conversation-alignment`. The repository is public and has 6 commits. The navigation tabs at the top are `<> Code`, `Issues`, `Pull requests`, `Actions`, `Projects`, `Security and quality`, and `Insights`. A green `<> Code` button is highlighted with a red box and the number 1. The file list on the left includes:


File	Description	Time
docs	Conversation alignment pipeline: wav -> word TextGrid -> sp...	1 hour ago
.gitattributes	Pin line endings: LF for shell scripts, CRLF for .bat	1 hour ago
.gitignore	Conversation alignment pipeline: wav -> word TextGrid -> sp...	1 hour ago
README.md	Conversation alignment pipeline: wav -> word TextGrid -> sp...	1 hour ago
TECHNICAL_NOTES.md	step 4: stop silently dropping words from the TextGrid	49 minutes ago
align_config.yaml	Conversation alignment pipeline: wav -> word TextGrid -> sp...	1 hour ago
check_setup.py	Fix g2p model resolution and fail fast on MFA 3.2+	1 hour ago
environment_mfa.yml	Conversation alignment pipeline: wav -> word TextGrid -> sp...	1 hour ago
environment_whisper.yml	Conversation alignment pipeline: wav -> word TextGrid -> sp...	1 hour ago
run_all.py	Add --no-g2p, and document that g2p may be costing us wo...	54 minutes ago
setup_mac.sh	Conversation alignment pipeline: wav -> word TextGrid -> sp...	1 hour ago
setup_windows.bat	Conversation alignment pipeline: wav -> word TextGrid -> sp...	1 hour ago
step1_whisper.py	UTF-8 stdout in every step script	1 hour ago
step2_mfa_corpus.py	UTF-8 stdout in every step script	1 hour ago

The right sidebar shows repository statistics: 0 stars, 0 forks, and 0 watchers. The language usage chart shows Python (85%), Shell (10.9%), and Batchfile (4.1%).

The page you should be looking at. The green **Code** button (1) is on the right, above the list of files — not the word “Code” in the row of tabs above it.

2 Download it

Click the green **Code** button, then **Download ZIP** at the bottom of the menu that opens. Ignore everything else on that menu — cloning, GitHub Desktop and the rest are for people editing the code.



The menu after clicking **Code**. You want **Download ZIP** (2), the last item. If your browser asks where to save it, Downloads is fine.

3 Unzip it into Documents

Open your Downloads folder, unzip the file (Windows: right-click → **Extract All**; Mac: double-click it), and move the resulting folder into **Documents**. Keep everything inside it together and do not rename the files.

The folder will be called `conversation-alignment-main`. That is normal — the `-main` is just how the download is named.

What is inside

FILE	WHAT YOU DO WITH IT
<code>setup_windows.bat</code> / <code>setup_mac.sh</code>	you run this once — it installs everything the code needs
<code>run_all.py</code>	you run this per recording — it makes the first-draft transcript
<code>textgrid_to_excel.py</code>	you run this at the end — turns your finished transcript into the spreadsheet
<code>check_setup.py</code>	tells you whether your computer is ready, in plain English
everything else	the machinery the three above use. You never open it

WINDOWS AND MAC WRITE PATHS DIFFERENTLY

You will never type a path by hand — the tricks in section 3 copy it for you — but it helps to recognise them:

WINDOWS

```
C:\Users\you\Documents\conversation-alignment-  
main
```

MAC

```
/Users/you/Documents/conversation-alignment-  
main
```

Back slashes on Windows, forward slashes on a Mac. Mixing them up is the single most common reason a command says it cannot find something.

2 Set up your computer — once ROUTE A ONLY

Two things to install, then one script that does the rest. It takes 30–45 minutes, nearly all of it waiting for downloads, and you never do it again. None of the lines below can break your computer.

1 Install Miniconda

Miniconda is Python plus a tool that stops different pieces of software fighting each other. Get it from docs.conda.io/en/latest/miniconda.html.

WINDOWS

Download the **Windows 64-bit** installer (`.exe`), run it, accept every default.

You now have a new Start-menu entry called **Anaconda Prompt**. That is the window you will work in.

MAC

You need the installer that matches your chip. Click the **Apple menu** at the top-left of the screen → **About This Mac**:

- **Apple M1 / M2 / M3...** → the **Apple Silicon (arm64)** `.pkg`
- **Intel** → the **Intel (x86_64)** `.pkg`

Double-click it, accept every default, then **quit Terminal completely (⌘+Q) and open it again** — until you do, your Mac will insist there is no such thing as `conda`.

2 Open the command window and go to the folder

WINDOWS

Start → type `Anaconda Prompt` → open it. **Not** Command Prompt or PowerShell; conda is only wired into this one.

Hold **Shift**, right-click the folder you unzipped, choose **Copy as path**. Then type `cd`, a space, and **right-click to paste** (Ctrl+V does nothing here):

```
cd "C:\Users\you\Documents\conversation-
alignment-main"
```

MAC

Press **⌘+Space**, type `Terminal`, press Enter.

Type `cd` and a space, then **drag the folder out of Finder onto the Terminal window** — it writes the path for you:

```
cd "/Users/you/Documents/conversation-
alignment-main"
```

You type one command, press **Enter**, and it prints text back. Nothing happens until you press Enter, and it never does anything you did not type.

3 Run the setup script

WINDOWS

```
setup_windows.bat
```

MAC

```
bash setup_mac.sh
```

That one line installs everything: the transcriber, the aligner, and the aligner's knowledge of English. It prints what it is doing as it goes and finishes by checking its own work.

It will sit apparently doing nothing for minutes at a time. That is the download, not a crash — these are several gigabytes. Leave the window alone until it stops.

When it is finished it prints **Ready**. If it prints anything else, it also prints the exact command to fix it; run that, then run the script again.

4 Let the code separate the speakers

Do this one if you can — it is what makes the speaker question in section 3 actually do something. Skip it and everything still runs, but every word lands on a single row and you separate the speakers by hand in Praat, which is a lot more clicking.

Make a free account at **huggingface.co**, open huggingface.co/pyannote/segmentation-3.0 and huggingface.co/pyannote/speaker-diarization-3.1, accept the terms on both, then run:

```
conda activate whisper
huggingface-cli login
```

It asks for a token, which you copy from your HuggingFace account settings. **The token stays invisible while you paste it** — that is deliberate. Press Enter anyway.

CHECK IT ANY TIME


```
python check_setup.py
```

It tests every piece and reports in plain English — something like:

```
[ok] conda      C:\Users\you\miniconda3\Scripts\conda.EXE
[ok] whisper env transcription libraries load
[ok] ffmpeg     present
[ok] mfa2 env    aligner version 3.1.0
[ok] language   english_us_arpa models downloaded
[ok] speakers   HuggingFace sign-in found
```

Ready. Transcribe a recording with: `python run_all.py`

Anything marked `X` comes with the command that fixes it. Run that, then check again.



```
Anaconda Prompt - call con... x + v
(whisper) C:\Users\you\Documents\conversation-alignment-main>python check_setup.py
Checking your setup. This takes about half a minute.

[ok] conda      C:\Users\james\anaconda3\Scripts\conda.EXE
[ok] whisper env transcription libraries load
[ok] ffmpeg     present
[ok] mfa2 env    aligner version 3.1.0
[ok] language   english_us_arpa models downloaded
[ok] speakers   HuggingFace sign-in found

Ready. Transcribe a recording with: python run_all.py

(whisper) C:\Users\you\Documents\conversation-alignment-main>
```

What that looks like on screen: a command window with every line **ok** and **Ready** at the bottom. This is the moment setup is finished.

WHEN SETUP MISBEHAVES

WHAT YOU SEE	WHAT IT MEANS
'conda' is not recognized (Windows)	you are in Command Prompt or PowerShell. Close it and open Anaconda Prompt instead
zsh: command not found: conda (Mac)	quit Terminal completely with ⌘+Q and open it again — the installer only changes windows opened after it ran
'setup_windows.bat' is not recognized	you are not in the folder. Do the <code>cd</code> in step 2 first, then <code>dir</code> (Windows) or <code>ls</code> (Mac) to check you can see the file
it sits there for minutes	normal — these downloads are gigabytes. Leave it alone
a download dies partway	run the setup script again; it carries on from where it stopped
CondaHTTPError	the network is blocking it — try again off the hospital wifi
red text but it finishes	usually harmless warnings. Run <code>python check_setup.py</code> ; if that says Ready, you are fine

3 Run the code on your recording ROUTE A ONLY

1 Get into the right place

```
conda activate whisper  
cd "the folder you unzipped"
```

Use the same `cd` trick as before — Shift+right-click → **Copy as path** on Windows, drag the folder onto Terminal on a Mac. Your prompt should now start with `(whisper)`.

2 Run it — it asks you three questions

```
python run_all.py
```

That is the whole command; you do not type any paths into it. It then asks you three things, one at a time:

IT ASKS	WHAT YOU DO
1 • the recording	drag the sound file into the window and press Enter. Dragging it in is the point — typing a path by hand is the most common way to get stuck here
2 • where the results go	press Enter to accept the folder it offers, which sits next to the recording. That is almost always what you want
3 • how many people are talking	type the number — 2, 3, 4 ... — if you know it. Press Enter only if you genuinely do not

It repeats your three answers back before it starts, so you can check them.

ANSWER THE SPEAKER QUESTION IF YOU POSSIBLY CAN

Counting the voices is the one thing the code is worst at — and the one thing you already know just by listening. Count **everyone whose voice is on the recording**: the person being recorded, whoever is asking the questions, and anyone else in the room who says more than a few words. Someone who only puts their head round the door does not count.

Give it a number and it splits the recording into **exactly** that many speaker rows. Press Enter instead and it guesses — and a wrong guess means one person split across two rows, or two people merged onto one, which you then untangle by hand in Praat. Two minutes of counting saves an hour of that.

Answer **1** if there is only one voice; it then skips the speaker step altogether and runs faster.

One catch: the split needs the HuggingFace sign-in from section 2, step 4. Without it every word lands on one row whatever you answer — the script says so before it starts, and you then separate the speakers yourself in Praat.

3 Leave it alone while it works

It reports as it goes, and runs the whole pipeline — transcription and alignment — without you doing anything else. Roughly what you should see:

```
[whisper] large-v3 on cuda (compute_type=float16)
  whisper @ 8.7 min, 1175 words
[done] 286 words, speakers ['SPEAKER_00']
[corpus] SPEAKER_00: 28 utterances
[mfa ] using the 'mfa2' environment - you do not need to switch to it yourself
[map ] 286/286 words placed; 22 kept Whisper's timing
```

WINDOWS WITH A GRAPHICS CARD

The first line says `on cuda`. A 30-minute recording takes a few minutes.

MAC, OR ANY MACHINE WITHOUT ONE

It says `on cpu (compute_type=int8)` instead. Nothing is wrong — it is simply slower. Budget roughly as long as the recording itself, sometimes more.

The first run also downloads the transcription model, around 3 GB, so it is far slower than every run after it.

4 Check you got a transcript

Your results folder should now hold `<name>_words.TextGrid`, and the last line should read `Done.` followed by the two files to open in Praat.

That TextGrid is your first draft. A `_words.xlsx` appears as well — ignore it; you will make the real spreadsheet at the end, after your corrections.

5 Next time, if you prefer

You can answer everything up front instead of being asked. The script prints this exact line for you at the end of the questions:

WINDOWS

```
python run_all.py "C:\path\to\recording.wav"
"C:\path\to\results" --speakers 3
```

MAC

```
python run_all.py "/path/to/recording.wav"
"/path/to/results" --speakers 3
```

IF IT STOPPED WITH AN ERROR

Run `python check_setup.py` first — most failures are a missing piece of setup, and it will name the one command that fixes it. If that says **Ready** and the run still fails, copy the **last ten lines** from the window and send them to whoever gave you the recording.

Do not start editing a draft the code did not finish. Its timings will be wrong throughout, and you would spend days fixing by hand what one re-run fixes in minutes.

EVERYONE READS FROM HERE ON

This is where **route B** starts. If a `.TextGrid` was made for you, nothing in sections 1–3 applies to you — you need Praat and nothing else.

4 Install Praat

Praat is free and needs no installation or admin rights — it is a single program you run.

1 Go to praat.org

Click **Download** for your computer (Windows or Mac).



Praat: doing phonetics by computer



Download Praat:

- * [Macintosh, Windows](#)
- * [Linux, Raspberry Pi, Chromebook](#)
- * [\(FreeBSD, SGI, Solaris, HP-UX\)](#)
- * [license](#) and [source code](#)

Information on Praat:

- * Introductory tutorial: choose **Intro** from Praat's **Help** menus.
- * Extensive manuals and tutorials: in Praat's **Help** menus.
- * [Beginner's manuals by others](#).
- * Paul Boersma's [publications](#) on algorithms and tutorials.

The authors:
Paul Boersma, David Weenink, Anastasia Shchupak
[Phonetic Sciences](#), University of Amsterdam
visiting: Spuistraat 134
mail: P.O. Box 1642, 1000BP Amsterdam
The Netherlands



[Paul](#)



[David](#)



Anastasia

Questions, problems, solutions:

1. Many problems can be solved by upgrading to [version 7.0.02](#) of Praat.
2. Make sure you have read the [Intro](#) from Praat's **Help** menu.
3. If that does not help, use the [Search](#) button in Praat's manual window.
4. Or consult the [Frequently Asked Questions](#) directly.
5. There is a user group on the Internet: the [Praat Users List](#).
6. If none of the above helps, you can send email to paul.boersma@uva.nl.

Functionality

The following gives you an idea of the features of the Praat program. The links take you into the web copy of the manual. The same manual is also available from Praat's Help menus, in which

praat.org — the download links are on the left.

2 Pick the right edition

On Windows take the one marked “**you probably want this**” (x64/v3). On a Mac, take the Apple-silicon edition unless your Mac is old.

Downloading Praat for Windows

1. Downloading your Windows edition

To download the latest version of the **Windows** (10, 11...) edition of **Praat**, download one of the following zip folders to your desktop:

x64(v3) edition (also known as Intel64, AMD64, x86_64...): [praat7002_win-x64v3.zip](#) (27 August 2026; 54.6 MB) <-- **you probably want this**
ARM64 edition: [praat7002_win-arm64.zip](#) (27 August 2026; 53.0 MB)
x64(v1) edition (also known as Intel64, AMD64, x86_64...): [praat7002_win-x64v1.zip](#) (27 August 2026; 54.2 MB)
32-bit edition (also known as Intel32, IA-32, x86, i386, i686...): [praat7002_win-intel32.zip](#) (27 August 2026; 55.1 MB)

The Windows download page. The top link is the one nearly everyone needs.

3 Unzip it and keep it somewhere sensible

The download is a zip with `Praat.exe` inside. Unzip it (right-click → Extract All) into a folder you will find again — Documents is fine. Double-click `Praat.exe` to start it.

Two windows open: **Praat Objects** (your file list) and **Praat Picture** (ignore it, you can close it).

5 Open your two files


You will be given two files: **the recording** (`.wav`) and **its transcript** (`.TextGrid`). They belong to the same session — keep them together.

1 Open → Read from file...

Pick the `.wav`. Then do it again and pick the `.TextGrid`. Both appear in the list.

2 Select BOTH, then click View & Edit

Click the first, then `Ctrl`-click the second so both are highlighted. The buttons on the right change depending on what is selected — with both selected, **View & Edit** opens them together.




1 both files highlighted — sound and transcript together. **2** the buttons appear in this panel; click **View & Edit**.

IF YOU ONLY SEE WORDS AND HEAR NOTHING

You opened the TextGrid on its own. Close it, select **both** files, and press View & Edit again.

6 What you are looking at



1 waveform — loudness over time. **2** spectrogram — the same sound by frequency; this is where you see a word begin. **3** the transcript: one row per speaker, one word per box. **4** the speaker names — click a name to work on that row. **5** a boundary — the line between two boxes, which you drag to move a word's start.

This is a real conversation from the lab, already cleaned. Empty boxes are silence — that is normal and correct.

7 How to actually edit

FIRST, ALWAYS

Click the **speaker name** (marker 4 above) for the row you want to change. Praat only edits the row you have selected — this is the number-one reason a key “does nothing”.

1 Listen

Drag across a stretch of the waveform to select it, then press **Tab** to play it. **Tab** again replays. Zoom in with **Ctrl** + **I** until you can see individual words.

2 Move a word's start — drag the boundary

Put the mouse on the vertical line at the start of the word, hold the left button, and drag it to where the sound really begins. Let go. That is the whole operation, and it is most of the job.

3 Fix the text — click the box and type

Click inside a word's box (it turns pink) and type. What you type replaces what is there. Click elsewhere or press **Tab** to move on.

4 Add a missing word — add a boundary

Click where the missing word should start, press **Enter** to drop a boundary, then click in the new box and type the word. Add a second boundary for where it ends if you need one.

Ctrl + **1** ... **8** does the same on a numbered row without selecting it first.

5 Get rid of a boundary you don't want

Click the boundary so it turns red (that selects it), then use the menu: **Boundary** → **Remove**. The two boxes either side merge into one, and their text is joined. There is no keyboard shortcut for this one.

6 Get rid of a word

Click its box and delete the text, so the box is empty. If the box itself should not be there, also remove the boundaries on each side (**Boundary** → **Remove**) so it merges back into the silence around it.

7 Save — often

File → **Save TextGrid as text file...**, keep the same name and folder. After the first time,

Ctrl + **S** saves. Save at the end of every work session.

Every key you need

Praat's own help writes “Command” and “Option”. On Windows those are **Ctrl** and **Alt**; on a Mac they are **⌘** and **⌥**.

WINDOWS	MAC	WHAT IT DOES
		play the selected interval — your most-used key. Press again to stop
+	+	zoom in
+	+	zoom out
+	+	zoom to what you have selected
+	+	zoom all the way back out
		add a boundary at the cursor, on the row you have selected
+ ...	+ ...	add a boundary on row 1...8 without selecting it first
+	+	save
+ / +	+ / +	undo / redo
+ / +	+ / +	find a word in the selected row / find the next one
drag a boundary		move a word's start or end
click a box		select that word — then type to change its text
Boundary → Remove		delete the selected boundary (menu only, no shortcut on either platform)

MOVING ALONG THE FILE

There is no keyboard shortcut for “next word” — click the next box, or drag a new selection and press . The arrow keys move the cursor inside the *text* you are typing, not between boxes.


8 What to type

Write what was actually said, not what should have been said. **Use these exact conventions** — the files get read by scripts, so consistency matters more than detail.

SITUATION	WHAT TO TYPE	EXAMPLE
a normal word	the word, one per box	dog
punctuation	keep it , attached to the word	Hello. right? Yeah,
“um”, “uh”, “mm”	type them — they are real and we want them	um uh
laughter	always exactly this, in its own box — not <i>hahaha</i> , not <i>heheh</i>	haha
you cannot make it out	always exactly this, in its own box. One word you cannot catch → one box. A mumbled stretch where you cannot even count the words → one box across the whole stretch , not a guess	??
cut-off word	what was said, plus a hyphen	wha-
repeats and stutters	keep every one	I I I think
numbers	as spoken	twenty , not 20
anything that is not a word: cough, laugh mid-word, beep, door, silence	leave the box empty	

ONE SPELLING EACH — DO NOT IMPROVISE

Older files are full of X , xxx , Hehehehh , EhehehehehHEH , Eheh — every one spelled differently, so no script can find them. **Do not copy that.** Laughter is always haha , however long the laugh is. Unclear speech is always ?? .



A filled pause gets its own box like any other word. The machine throws these away, so they are usually missing until you add them.

ONE WORD, ONE BOX, ONE SPEAKER

If two people say the same thing at once, it goes on **one** row only — whoever actually said it. Putting it on both rows counts that word twice in every later analysis.

A VOICE THAT IS NOBODY ON THE LIST

Nurses, techs and visitors wander in. How much they say decides what you do:

- **Just a few words** (a greeting, a set of vitals — say under ~10 words): put them on a row named `other`, added once at the bottom. Not worth more.
- **They actually take part** in the conversation: give them their own row, named for who they are, the same as any other speaker — they are a real participant and the analysis needs them separated.

Either way, never park a stranger's speech on someone else's row: in one recording a nurse's vitals ended up on a researcher's row, which quietly corrupts that speaker's data.


9 Where a word starts

This is what the science depends on. A word's **start** matters far more than its end.


- **Vowels and m/n/l/r:** where the regular striping begins in the spectrogram — voicing starting up.

- **p, t, k, b, d, g:** at the **burst** (the sharp vertical spike), not the silence before it.
- **s, f, sh:** where the fuzzy high-frequency noise starts.
- **After a pause:** where the sound first lifts off the flat line.

Aim to be within about **20 ms**. On this room-microphone audio, 20–30 ms is realistic and good.



A correctly placed boundary: right where the sound for that word begins.




Why you check every one. Same words, two versions: the top row is the computer's raw guess, the bottom is corrected. The machine's boundaries sit late, consistently.

IF EVERY BOUNDARY LOOKS LATE IN THE SAME WAY

Stop and ask the person who ran Whisper — the file probably missed its alignment step, which is quick to redo.

Two people talking at once



Overlapping speech: both voices are in the same waveform, so no boundary here can be exact.

Get the words right, accept looser boundaries, and move on. Roughly 3–20% of a conversation is crosstalk and it is nobody's fault.

WHEN YOU GENUINELY CANNOT TELL

Leave the boundary where it is and note the timestamp. An honest “not checked” is worth more than a confident guess.

10 How to work through a file

1 Listen through once without editing

Ten minutes at speed, to learn the voices and who is in the room. Everything after this goes faster.

2 Fix words first, boundaries second

A perfectly placed boundary on the wrong word is wasted work.

3 Work in 10–15 minute blocks of audio

Save at the end of each block. A 30-minute recording is a 6–10 hour job — not one sitting.

4 Work on a copy, and keep the rows as they are

Before you start, copy the TextGrid and add your initials and the date: `<original name>_ab20260924.TextGrid`. Edit that copy and leave the original untouched — whoever takes it next renames it to the final name.

Inside the file, do not rename rows or add and delete them, with the one exception of an `other` row for a voice that belongs to nobody.

5 Write down where you stopped

If you did not finish, say exactly where you got to. A file corrected for ten minutes and raw after that, with nothing saying so, is worse than one nobody touched.

11 Turn it into the spreadsheet

The lab works from an Excel file, not the Praat file. One command makes it from your finished transcript:

```
conda activate whisper
python textgrid_to_excel.py "your_file.TextGrid" "your_file.xlsx" --exclude
```

You get one row per word: `onset`, `offset` and `Duration` in milliseconds, then one column per speaker with the word in that speaker's column. `--exclude` adds two more columns marking the rows that are not words (`??` and `haha`) so they can be filtered out in one step.

Hand back **both** files — the `.TextGrid` and the `.xlsx`.

ON ROUTE B, AND NEVER INSTALLED ANY OF THIS?

Then you do not have the `whisper` environment that command needs. Either **hand back just your `.TextGrid`** and say you have not made the spreadsheet — whoever gave you the file can do it in one command — or install the small part on its own, which takes a couple of minutes and none of the large downloads in section 2:

```
conda create -n convo python=3.11 -y
conda activate convo
pip install pandas praatio openpyxl
```

Then run the `textgrid_to_excel.py` line above with `convo` in place of `whisper`.

12 Before you hand it back

- ☐ every word matches what is said, punctuation kept

- ☐ “um” and “uh” typed in
- ☐ ?? and haha used — no X, no invented spellings
- ☐ working on a copy with your initials and date, original untouched
- ☐ one word per box; no word sitting on two speaker rows
- ☐ word starts checked against the waveform, not left where the machine put them
- ☐ no overlapping boxes on the same row, no zero-width boxes
- ☐ file name, row names and folder unchanged
- ☐ saved as a text file, and the spreadsheet made with `textgrid_to_excel.py`
- ☐ if you did not finish, say where you stopped when you hand it back

STUCK?

Ask whoever assigned you the file. A five-minute question is cheaper than a day spent on the wrong one, and “this file looks wrong” has been the right call more than once.

Screenshots are real Praat windows on an already-cleaned lab recording.